



**Michigan
Technological
University**

MICHIGAN TECHNOLOGICAL UNIVERSITY, COMPUTER SCIENCE

MA 5790 Combined Section - Predictive Modeling Assignment 3

Anil Silwal

November 2, 2018

Problem 6.1

- Start R and use these commands to load the data
- In this example the predictors are the measurements at the individual frequencies. Because the frequencies lie in a systematic order ($850 \sim 1,050 \text{ nm}$), the predictors have a high degree of correlation. Hence, the data lie in a smaller dimension than the total number of predictors (100). Use PCA to determine the effective dimension of these data. What is the effective dimension.
- Split the data into a training and a test set, pre-process the data, and build each variety of models described in this chapter. For those models with tuning parameters, what are the optimal values of the tuning parameter(s)?
- Which model has the best predictive ability? Is any model significantly better or worse than the others?
- Explain which model you would use for predicting the fat content of a sample.




Fig. 6.20: A sample of ten spectra of the Tecator data. The colors of the curves reflect the absorption values, where *yellow* indicates low absorption and *red* is indicative of high absorption

Solution 6.1(b):

Applying the PCA, the results shows that the first component which has PCA value = 98.6261, contains almost all information. This means the true dimensionality is much lower than number of predictors.

In principal component analysis(PCA), we often use the scree plot to figure out how many important components to include in the model. Plotting the screeplot using "screeplot()" function, the elbow comes at the 2nd component in our data and the first two components in our model would be best to retained. **Thus, the first 2 component would be the effective dimension.**

However, this is based on linear combination of data; the other non-linear summarization of data may be also useful.




Figure 1: Scree Plot showing important Components

Solution 6.1(c, d, e):

Given, we have 3 endpoints - water, protein and fat. And as per the given question, the predictive relationship between IR spectrum and fat content is required to provide cost saving as analytical chemistry is a more expensive, time-consuming process.

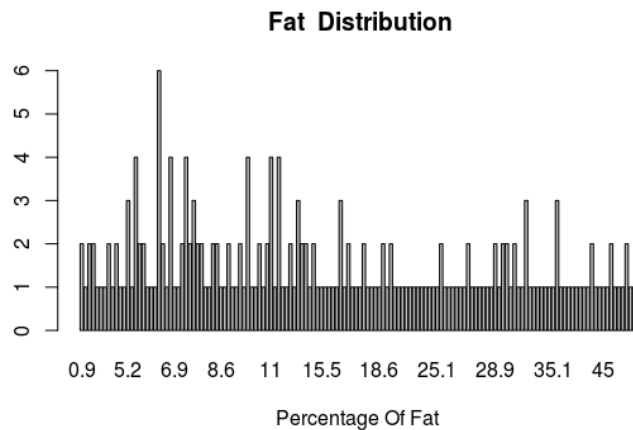


Figure 2: Response Predictor Distribution i.e Fat

Given the sample size, we split the data into 80% and 20% and repeat = 4 10-fold cross validation to tune our model. For resampling, I tried $k=4$, $k=8$ and $k=10$. At last the 10 fold cross-validation resampling gives better RMSE value and I choose the 10 folded resampling and repeated 4 times in a given simpler linear regression method.

```

Linear Regression

163 samples
100 predictors

No pre-processing
Resampling: Cross-Validated (4 fold)
Summary of sample sizes: 122, 122, 123, 122
Resampling results:

    RMSE      Rsquared   MAE
7.097009  0.750529  3.704548

Tuning parameter 'intercept' was held constant at a value of TRUE

```

Figure 3: Resampling for $k = 4$

```

Linear Regression

163 samples
100 predictors

No pre-processing
Resampling: Cross-Validated (8 fold)
Summary of sample sizes: 142, 142, 143, 144, 143, 143, ...
Resampling results:

    RMSE      Rsquared   MAE
4.087294  0.9102864  2.350725

Tuning parameter 'intercept' was held constant at a value of TRUE

```

Figure 4: Resampling for $k = 8$

```

Linear Regression

163 samples
100 predictors

No pre-processing
Resampling: Cross-Validated (10 fold, repeated 4 times)
Summary of sample sizes: 146, 147, 146, 147, 147, 147, ...
Resampling results:

    RMSE      Rsquared   MAE
3.981347  0.9035059  2.29246

Tuning parameter 'intercept' was held constant at a value of TRUE

```

Figure 5: Resampling for $k = 10$

The prediction accuracy based on this simple linear regression model is:

RMSE	Rsquared	MAE
3.7758443	0.9272617	2.3082520




Figure 6: Simple Linear Regression Model - Prediction's Accuracy

After that, I applied PCR and PLS method to the given data:

Principal Component Analysis

163 samples
100 predictors

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 4 times)

Summary of sample sizes: 146, 146, 146, 147, 146, 147, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	MAE
1	11.285269	0.2375411	9.175386
2	11.268372	0.2585440	9.035455
3	8.495126	0.5685672	6.658762
4	4.298917	0.8981045	3.559623
5	3.388837	0.9464940	2.695047
6	3.099499	0.9546983	2.460895
7	3.090544	0.9545041	2.439653
8	3.087446	0.9566478	2.399654
9	2.964698	0.9587683	2.345524
10	2.978828	0.9580843	2.283250
11	2.757382	0.9638273	2.196121
12	2.764683	0.9633439	2.205798
13	2.805820	0.9622680	2.230794
14	2.871945	0.9609657	2.263338
15	2.820530	0.9631437	2.188388
16	2.610498	0.9678740	1.979684
17	2.586934	0.9678644	1.955452
18	2.528137	0.9685569	1.930185
19	2.597241	0.9655178	1.949137
20	2.712400	0.9608449	1.999735
21	2.643882	0.9622737	1.910566
22	2.481492	0.9667971	1.819169
23	2.518659	0.9658297	1.834688
24	2.568508	0.9643670	1.848284

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was ncomp = 22.




Figure 7: PCR Model

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
2.676961	0.962550	1.953781




Figure 8: PCA Model - Prediction's Accuracy

Partial Least Squares

163 samples
100 predictors

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 4 times)

Summary of sample sizes: 146, 147, 146, 146, 148, 147, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	MAE
1	11.210142	0.2533495	9.135502
2	7.061965	0.6969594	5.450560
3	5.219216	0.8242954	4.039968
4	4.142995	0.9002930	3.431444
5	3.135929	0.9501175	2.422840
6	3.045431	0.9519150	2.401466
7	3.003821	0.9544468	2.316960
8	2.925475	0.9564938	2.240443
9	2.823914	0.9567421	2.216611
10	2.737971	0.9592964	2.156776
11	2.699708	0.9608291	2.164580
12	2.603303	0.9637960	1.981274
13	2.379695	0.9702258	1.809272
14	2.383486	0.9694343	1.772381
15	2.606049	0.9614942	1.880854
16	2.772072	0.9527932	1.891770
17	2.666571	0.9569108	1.840283
18	2.559881	0.9585328	1.760982
19	2.488756	0.9596306	1.707179
20	2.607386	0.9526803	1.731890
21	2.878102	0.9348846	1.796906
22	3.088140	0.9160315	1.827737
23	3.109935	0.9120810	1.831501
24	3.087757	0.9125988	1.835312

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was ncomp = 13.




Figure 9: PLS Model

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
2.395839	0.966689	1.826058




Figure 10: PLS Model - Prediction's Accuracy

Ridge Regression

163 samples
100 predictors

Pre-processing: centered (100), scaled (100)

Resampling: Cross-Validated (10 fold, repeated 4 times)

Summary of sample sizes: 146, 147, 146, 146, 147, 147, ...

Resampling results across tuning parameters:

lambda	RMSE	Rsquared	MAE
0.00000000	3.625519	0.9229186	2.276718
0.07142857	4.802791	0.8742580	3.793913
0.14285714	5.586271	0.8308913	4.349821
0.21428571	6.171174	0.7887722	4.728845
0.28571429	6.635387	0.7493139	5.040823
0.35714286	7.024822	0.7131838	5.302266
0.42857143	7.366853	0.6805018	5.535095
0.50000000	7.678371	0.6511182	5.751513
0.57142857	7.970192	0.6247635	5.956757
0.64285714	8.249463	0.6011284	6.157070
0.71428571	8.521023	0.5799047	6.363167
0.78571429	8.788208	0.5608050	6.577567
0.85714286	9.053341	0.5435702	6.797944
0.92857143	9.318055	0.5279717	7.021094
1.00000000	9.583493	0.5138100	7.252147

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was lambda = 0.




Figure 11: Ridge Regression - Model

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
3.7758856	0.9272604	2.3082759




Figure 12: Ridge Regression - Prediction's Accuracy

The lasso

163 samples

100 predictors

Pre-processing: centered (100), scaled (100)

Resampling: Cross-Validated (10 fold, repeated 4 times)

Summary of sample sizes: 147, 146, 147, 146, 146, 147, ...

Resampling results across tuning parameters:

fraction	RMSE	Rsquared	MAE
0.1000000	2.557048	0.9586631	1.663415
0.1473684	2.536936	0.9569477	1.620058
0.1947368	2.489033	0.9581858	1.611681
0.2421053	2.504299	0.9565958	1.612422
0.2894737	2.588675	0.9537632	1.643274
0.3368421	2.665490	0.9511100	1.676271
0.3842105	2.754083	0.9477797	1.714502
0.4315789	2.843932	0.9442388	1.765119
0.4789474	2.940890	0.9402209	1.818507
0.5263158	3.042442	0.9359181	1.877187
0.5736842	3.161280	0.9307773	1.945045
0.6210526	3.294699	0.9246004	2.010607
0.6684211	3.430530	0.9185784	2.078964
0.7157895	3.554727	0.9133475	2.142798
0.7631579	3.689140	0.9077066	2.208040
0.8105263	3.826560	0.9020930	2.278452
0.8578947	3.961111	0.8968452	2.349657
0.9052632	4.096749	0.8916446	2.419650
0.9526316	4.238345	0.8863030	2.491619
1.0000000	4.375948	0.8811980	2.562891

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was fraction = 0.1947368.




Figure 13: Lasso - Model

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
3.7779787	0.9205828	2.0205018




Figure 14: Lasso Model - Prediction's Accuracy

Elasticnet

163 samples
100 predictors

Pre-processing: centered (100), scaled (100)
Resampling: Cross-Validated (10 fold, repeated 4 times)
Summary of sample sizes: 147, 146, 147, 147, 147, 147, ...
Resampling results across tuning parameters:

lambda	fraction	RMSE	Rsquared	MAE
0.000	0.05	2.296717	0.9692331	1.599045
0.000	0.10	2.544109	0.9582517	1.669316
0.000	0.15	2.430343	0.9613246	1.601947
0.000	0.20	2.358950	0.9640401	1.568385
0.000	0.25	2.349822	0.9654756	1.561296
0.000	0.30	2.366620	0.9657882	1.561860
0.000	0.35	2.419151	0.9638669	1.587187
0.000	0.40	2.495391	0.9606014	1.626680
0.000	0.45	2.580223	0.9572505	1.670147
0.000	0.50	2.682988	0.9534439	1.728129
0.000	0.55	2.812882	0.9487650	1.794103
0.000	0.60	2.953915	0.9432250	1.864652
0.000	0.65	3.082659	0.9380816	1.928348
0.000	0.70	3.203102	0.9331208	1.990805
0.000	0.75	3.325725	0.9279904	2.053661
0.000	0.80	3.451731	0.9230643	2.117962
0.000	0.85	3.576532	0.9182862	2.185276
0.000	0.90	3.708254	0.9131419	2.256037
0.000	0.95	3.839454	0.9079909	2.324038
0.000	1.00	3.982111	0.9023055	2.395655
0.001	0.05	9.967684	0.4433099	8.139487
0.001	0.10	8.689417	0.6227025	7.121237
0.001	0.15	7.481911	0.7476191	6.145428
0.001	0.20	6.383120	0.8235283	5.198319

Figure 15: Elastic Net Model

0.001	0.20	6.383120	0.8235283	5.198319
0.001	0.25	5.512517	0.8654279	4.431189
0.001	0.30	4.761669	0.8968458	3.822577
0.001	0.35	4.143522	0.9185019	3.372292
0.001	0.40	3.730624	0.9316744	3.105014
0.001	0.45	3.506135	0.9388213	2.928610
0.001	0.50	3.386530	0.9430143	2.812688
0.001	0.55	3.301222	0.9464624	2.716169
0.001	0.60	3.227780	0.9494004	2.626238
0.001	0.65	3.179387	0.9513161	2.559868
0.001	0.70	3.147200	0.9525621	2.513164
0.001	0.75	3.125075	0.9533745	2.480941
0.001	0.80	3.109820	0.9539287	2.458055
0.001	0.85	3.095848	0.9543971	2.438064
0.001	0.90	3.084022	0.9547755	2.421701
0.001	0.95	3.074917	0.9550582	2.409538
0.001	1.00	3.068482	0.9552464	2.402532
0.010	0.05	10.519126	0.3559662	8.596385
0.010	0.10	9.753087	0.4747515	7.961266
0.010	0.15	9.011085	0.5797539	7.372586
0.010	0.20	8.292520	0.6666571	6.805026
0.010	0.25	7.602207	0.7349230	6.243681
0.010	0.30	6.963250	0.7851325	5.702748
0.010	0.35	6.375113	0.8213783	5.188552
0.010	0.40	5.876665	0.8464846	4.746301
0.010	0.45	5.436747	0.8670412	4.382013
0.010	0.50	5.014000	0.8851875	4.039269
0.010	0.55	4.664669	0.8983297	3.752973
0.010	0.60	4.369346	0.9083816	3.537720
0.010	0.65	4.146085	0.9152606	3.402399
0.010	0.70	3.996934	0.9194388	3.308889
0.010	0.75	3.886617	0.9226141	3.229740
0.010	0.80	3.807617	0.9250116	3.170508
0.010	0.85	3.750909	0.9268554	3.125558
- - - -	- - -	- - - - -	- - - - -	- - - - -

Figure 16: Elastic Net Model- Contd.

0.100	0.90	5.312216	0.8498342	4.201370
0.100	0.95	5.209910	0.8529736	4.115605
0.100	1.00	5.122362	0.8554433	4.038504
1.000	0.05	11.205401	0.3045177	9.440021
1.000	0.10	10.832274	0.2991451	8.674205
1.000	0.15	11.434928	0.2966806	8.875386
1.000	0.20	11.639567	0.3088883	9.043382
1.000	0.25	11.450751	0.3274330	8.879424
1.000	0.30	11.270876	0.3453515	8.718568
1.000	0.35	11.097908	0.3627884	8.561847
1.000	0.40	10.932342	0.3797007	8.414843
1.000	0.45	10.775017	0.3959795	8.279911
1.000	0.50	10.624672	0.4117068	8.151221
1.000	0.55	10.484699	0.4265969	8.035107
1.000	0.60	10.352857	0.4407918	7.926113
1.000	0.65	10.231595	0.4539749	7.822487
1.000	0.70	10.126556	0.4654350	7.730143
1.000	0.75	10.030389	0.4755232	7.646175
1.000	0.80	9.935486	0.4851168	7.565204
1.000	0.85	9.842909	0.4943127	7.487894
1.000	0.90	9.754275	0.5030871	7.413591
1.000	0.95	9.670815	0.5113741	7.343148
1.000	1.00	9.593016	0.5191247	7.277101

RMSE was used to select the optimal model using the smallest value.
The final values used for the model were fraction = 0.05 and lambda = 0.

Figure 17: Elastic Net Model- Contd.




Figure 18: Elastic Net Graph

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
2.4802401	0.9673689	1.7657605




Figure 19: Elastic Net Model- Prediction

Solution 6.1(d)

Comparing the results from different models, The **PLS** has minimum RMSE which is 2.395839. The **PCA** and **Elastic Net** also has nearly similar RMSE which are 2.676961 and 2.4802401 respectively. This suggests, **PLS**, **PCA**, **Elastic Net** are better than others.

While, the **Lasso** has maximum RMSE which is 3.77797. And this Lasso Model is comparatively worst model for this dataset. Other than that, **Simple Regression** Model has also worst performance and shows, RMSE = 3.7758856 which is also the worst case in comparison to other models.

Solution 6.1(e)

Comparing the results from different models, the **PLS** also has minimum RMSE which is 2.395839. This suggests, **PLS** is most likely used model for predicting Fat in this dataset.

Comparison table for different Performance

Model	RMSE Training	R^2 Training	RMSE Testing	R^2 Testing
Simple	3.981347	0.9035059	3.7758443	0.9272617
PCA	2.481492	0.9667971	2.676961	0.962550
PLS	2.379695	0.97.2258	2.395839	0.966689
Ridge	3.625519	3.625519	3.7758856	0.9272604
Lasso	2.489033	0.9581858	3.7779787	0.9205828
Elastic Net	2.296717	0.9692331	2.4802401	0.9673689

Problem 6.2 Developing a model to predict permeability (see Sect. 1.4) could save significant resources for a pharmaceutical company, while at the same time more rapidly identifying molecules that have a sufficient permeability to become a drug:

- Start R and use these commands to load the data
- The fingerprint predictors indicate the presence or absence of substructures of a molecule and are often sparse meaning that relatively few of the molecules contain each substructure. Filter out the predictors that have low frequencies using the *nearZeroVar* function from the caret package. How many predictors are left for modeling?
- Split the data into a training and a test set, pre-process the data, and tune a PLS model. How many latent variables are optimal and what is the corresponding resampled estimate of R^2 ?
- Predict the response for the test set. What is the test set estimate of R^2 ?
- Try building other models discussed in this chapter. Do any have better predictive performance
- Would you recommend any of your models to replace the permeability laboratory experiment?

Solution 6.2(b)

Initially, we have 1107 predictors in fingerprints. Applying non-zero-variance, there were 719 predictors found having non-zero-variance fingerprints. Then after removing these predictors, only 388 predictors are left for the modeling. The fingerprints predictors are reduced to almost 25 percent comparing to original predictors.

```
Before Non-Zero Variance, number of predictors in fingerprints is 1107:
num [1:165, 1:1107] 0 0 0 0 0 0 0 0 0 0 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:165] "1" "2" "3" "4" ...
..$ : chr [1:1107] "X1" "X2" "X3" "X4" ...
NULL

After Non-Zero Variance, number of predictors in fingerprints is 388:
num [1:165, 1:388] 0 0 0 0 0 0 0 0 0 0 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:165] "1" "2" "3" "4" ...
..$ : chr [1:388] "X1" "X2" "X3" "X4" ...
NULL
```

Figure 20: Result showing predictors before and after applying near-zero-variance() function

Solution 6.2(c)

When determining the data splitting method, we should focus on two primary characteristics:

- the number of samples relative to the number of predictors in the data, and
- the distribution of samples across classes.

After the near-zero-variance, the data sample i.e sample = 165 and predictors = 388. Still number of predictors are larger than the sample data. In such case, we use stratified random splitting and for that we use `createDataPartition()` function given by library(caret). For this, we used 75% as training data and 25% as testing data from sample. For PLS Tuning Parameter, the 4 fold cross validation with 5 repeats technique was used to

to resample number of train datas. For PLS tuning, I take 15 components and calculate the RMSE, Rsquared and MAE values.

Partial Least Squares

125 samples

388 predictors

Pre-processing: centered (388), scaled (388)

Resampling: Cross-Validated (4 fold, repeated 5 times)

Summary of sample sizes: 93, 93, 96, 93, 94, 94, ...

Resampling results across tuning parameters:

ncomp	RMSE	Rsquared	MAE
1	13.31396	0.2325657	9.729249
2	12.68750	0.3255819	9.073525
3	12.70913	0.3463417	9.351738
4	13.00558	0.3419348	9.482017
5	12.75070	0.3749557	9.167606
6	12.74253	0.3683852	9.258273
7	12.61378	0.3761673	9.251668
8	12.64845	0.3782935	9.442780
9	12.79393	0.3776872	9.520957
10	13.00754	0.3696524	9.661983
11	13.13951	0.3661599	9.658102
12	13.12691	0.3690678	9.658389
13	13.24511	0.3637436	9.759184
14	13.45138	0.3548575	9.843516
15	13.67831	0.3428259	10.006368

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was ncomp = 7.




Figure 21: Result showing PLS Tuning Parameter For Permeability Data

The result shows the optimal number of latent variables that maximizes R^2 is 7th component, which is 0.3761673.

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
10.8629421	0.5957486	8.4115944




Figure 22: PLS : Prediction-Accuracy

The result shows the corresponding resampled estimate of R^2 is 0.5957486, given by PLS.

Ridge Regression

125 samples
388 predictors

Pre-processing: centered (388), scaled (388)

Resampling: Cross-Validated (4 fold, repeated 5 times)

Summary of sample sizes: 93, 93, 94, 95, 94, 94, ...

Resampling results across tuning parameters:

lambda	RMSE	Rsquared	MAE
0.08000000	13.67989	0.3494300	9.952184
0.09214286	13.55044	0.3559901	9.863265
0.10428571	13.45758	0.3615510	9.803490
0.11642857	13.43308	0.3644957	9.803520
0.12857143	13.35137	0.3694529	9.749217
0.14071429	13.32510	0.3718913	9.739208
0.15285714	13.30119	0.3747566	9.730755
0.16500000	13.29086	0.3773629	9.727753
0.17714286	13.28210	0.3794440	9.732622
0.18928571	13.27738	0.3812073	9.744017
0.20142857	13.27802	0.3828925	9.753622
0.21357143	13.29302	0.3842816	9.778467
0.22571429	13.30597	0.3855011	9.802696
0.23785714	13.32910	0.3865729	9.837620
0.25000000	13.35203	0.3875474	9.869504

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was lambda = 0.1892857.




Figure 23: Ridge Regression - Model

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
10.4881473	0.6284279	8.3010729




Figure 24: Ridge Regression - Prediction's Accuracy

Elasticnet

125 samples
388 predictors

No pre-processing

Resampling: Cross-Validated (4 fold, repeated 5 times)

Summary of sample sizes: 93, 93, 96, 93, 93, 94, ...

Resampling results across tuning parameters:

lambda	fraction	RMSE	Rsquared	MAE
0.001	0.05	2248.30439	0.3028617	1431.826331
0.001	0.10	4344.72069	0.3300379	2595.133575
0.001	0.15	6237.23489	0.3310081	3697.777365
0.001	0.20	8053.96982	0.3105826	4773.030612
0.001	0.25	9894.57727	0.2926703	5895.466133
0.001	0.30	11620.14660	0.2825662	6944.799592
0.001	0.35	13325.46382	0.2718992	7983.205697
0.001	0.40	15037.32757	0.2590371	9022.937085
0.001	0.45	16747.97054	0.2394469	10049.994640
0.001	0.50	18447.65705	0.2240361	11038.085071
0.001	0.55	20159.48171	0.2086909	12068.071488
0.001	0.60	21878.45694	0.1952575	13101.193485
0.001	0.65	23551.82611	0.1829680	14124.846686
0.001	0.70	25222.70984	0.1715138	15145.877003
0.001	0.75	26898.48857	0.1625998	16166.612156
0.001	0.80	28578.38317	0.1558386	17187.095642
0.001	0.85	30280.53189	0.1510414	18212.079420
0.001	0.90	31977.36762	0.1482362	19234.080122
0.001	0.95	33666.74789	0.1466034	20252.692172
0.001	1.00	35358.06121	0.1453706	21271.209132
0.010	0.05	24.64205	0.3459478	17.038895
0.010	0.10	37.50072	0.3679587	25.773433
0.010	0.15	49.52693	0.3775678	33.939236
0.010	0.20	61.23008	0.3787936	41.711423
0.010	0.25	73.01430	0.3733106	49.426559
0.010	0.30	84.71737	0.3667700	57.036169
0.010	0.35	96.55428	0.3553505	64.757010

Figure 25: Elastic Net Model

0.010	0.35	96.55428	0.3553505	64.757010
0.010	0.40	108.44026	0.3446569	72.616348
0.010	0.45	120.33356	0.3345303	80.496988
0.010	0.50	132.12039	0.3264667	88.334571
0.010	0.55	143.90822	0.3178218	96.233371
0.010	0.60	155.49844	0.3092702	103.993373
0.010	0.65	166.52027	0.2994670	111.444983
0.010	0.70	177.51670	0.2907953	118.894999
0.010	0.75	188.55919	0.2816290	126.340504
0.010	0.80	199.66706	0.2741804	133.793491
0.010	0.85	210.72149	0.2674920	141.193621
0.010	0.90	221.47537	0.2618222	148.414469
0.010	0.95	232.20570	0.2564006	155.597899
0.010	1.00	242.91633	0.2507635	162.738791
0.100	0.05	12.88850	0.3432112	9.642427
0.100	0.10	12.41287	0.3548075	8.775103
0.100	0.15	12.29873	0.3726547	8.640358
0.100	0.20	12.20814	0.3876343	8.675209
0.100	0.25	12.12783	0.3982901	8.668984
0.100	0.30	12.12502	0.4011206	8.657021
0.100	0.35	12.18389	0.4001623	8.713620
0.100	0.40	12.25657	0.3994417	8.766410
0.100	0.45	12.36312	0.3969573	8.831905
0.100	0.50	12.48020	0.3943788	8.914045
0.100	0.55	12.62470	0.3905697	9.022619
0.100	0.60	12.74594	0.3882574	9.126388
0.100	0.65	12.85874	0.3861938	9.224724
0.100	0.70	12.97076	0.3842770	9.312166
0.100	0.75	13.09570	0.3818052	9.414117
0.100	0.80	13.21788	0.3793951	9.520797
0.100	0.85	13.32799	0.3774930	9.615461
0.100	0.90	13.43510	0.3745144	9.699796
0.100	0.95	13.54478	0.3713297	9.779136
0.100	1.00	13.66063	0.3678376	9.861984

RMSE was used to select the optimal model using the smallest value.
The final values used for the model were fraction = 0.3 and lambda = 0.1.

Figure 26: Elastic Net Model- Contd.

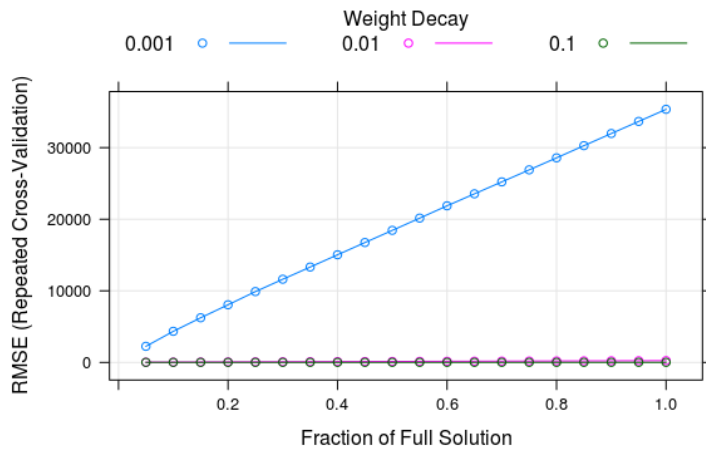


Figure 27: Elastic Net Graph

The prediction accuracy based on this model is:

RMSE	Rsquared	MAE
10.8629421	0.5957486	8.4115944




Figure 28: Elastic Net Model- Prediction

Comparing the results from different models, The **Elastic net** has minimum RMSE which is 10.0757239.

Solution 6.1(f)

Comparing the results from different models, the **Elastic net** has minimum RMSE which is 10.0757239. This suggests, **Elastic Net** is most likely used model for predicting Permeability in this dataset.

Comparison table for different Performance

Model	RMSE Training	R^2 Training	RMSE Testing	R^2 Testing
PLS	12.61378	0.3761673	10.8629421	0.5957486
Ridge	13.277328	0.3812073	10.4881743	0.6284279
Elastic Net	12.12502	0.4011406	10.0757239	0.6605789

```

library(mlbench)
library(caret)
library(e1071)
library(AppliedPredictiveModeling)
library(dplyr)
library(data.table)

#####
#Question 6.1
#####
library(AppliedPredictiveModeling)
library(MASS)
library(caret)
library(elasticnet)
library(lars)
library(pls)

data(tecator)

#####
# question 6.1(b)
#####

colName = -
for (i in 1:100)-
  colName[i] = paste("X",i)

colnames(absorp) = colName

## The base R function prcomp can be used for PCA. In the code below,
## the data are centered and scaled prior to PCA.
pcaObject = prcomp(absorp, center = TRUE, scale. = TRUE)

# The standard deviations for the columns in the data are stored in pcaObject as a sub-object called sd:
# Calculate the cumulative percentage of variance which each component
# accounts for.
percentVariance = pcaObject$sdev^2 / sum(pcaObject$sdev^2) * 100
print(percentVariance[1:5])

# taking only 5 variances set npcs = 5
screeplot(pcaObject, npcs = 5, type = "lines", main = "Scree Plot for PCA Analysis")

#####
#####

#####
# question 6.1(c,d,e)
#####

# Response variable = Fat

# BarPlot of response variable

counts = table(endpoints[,2])
barplot(counts, main="Fat Distribution",
        xlab="Percentage Of Fat ")

# splitting data into 80% and 20% based on Fat Response
set.seed(12345)

trainingRows = createDataPartition(endpoints[,2], p = .75, list= FALSE)

trainAbsorption = absorp[ trainingRows, ]

```

```

testAbsorption <- absorp[-trainingRows, ]
trainFat <- endpoints[trainingRows, 2]
testFat <- endpoints[-trainingRows, 2]

ctrl <- trainControl(method = "repeatedcv", repeats=4)

cat("n")

set.seed(12345)

# simple linear regression
meatlm <- train(x = trainAbsorption , y = trainFat, method = "lm", trControl = ctrl)
print(meatlm)

prediction<-predict(meatlm,testAbsorption)
accuracy3<-data.frame(obs=testFat,pred=prediction)
defaultSummary(accuracy3)
plot(accuracy3)

cat("n")

# PCR method
meatPCR <- train(x = trainAbsorption , y = trainFat, method = "pcr", trControl = ctrl, tuneLength = 24)
print(meatPCR)
plot(meatPCR)

cat("n")
cat("n")

# PLS method
meatPLS <- train(x = trainAbsorption , y = trainFat, method = "pls", trControl = ctrl, tuneLength = 24)
print(meatPLS)
plot(meatPLS)
cat("n")

# Ridge Regression Method
meatRg <- train(x = trainAbsorption , y = trainFat, method = "ridge",
               trControl = ctrl,
               preProcess = c("center","scale"),
               tuneGrid = expand.grid(lambda = seq(0,1,length=15)))

print(meatRg)
plot(meatRg)
cat("n")

prediction<-predict(meatRg,testAbsorption)
accuracy2<-data.frame(obs=testFat,pred=prediction)
defaultSummary(accuracy2)
plot(accuracy2)

# Lasso Regression Method
meatLasso <- train(x = trainAbsorption , y = trainFat, method = "lasso",
                  trControl = ctrl,
                  preProcess = c("center","scale"),
                  tuneGrid = expand.grid(fraction = seq(0.1,1,length=20)))

print(meatLasso)
plot(meatLasso)
cat("n")

```

```

prediction;-predict(meatLasso,testAbsorption)
accuracy3;-data.frame(obs=testFat,pred=prediction)
defaultSummary(accuracy3)
plot(accuracy3)

```

```

# Elastic Net Method
meatEls ; train(x = trainAbsorption , y = trainFat, method = "enet",
               trControl = ctrl,
               preProcess = c("center","scale"),
               tuneGrid = expand.grid(lambda = c(0,.001,.01,.1,1),
                                     fraction = seq(0.05,1,length=20)))

```

```

print(meatEls)
plot(meatEls)
cat("n")

```

```

prediction;-predict(meatEls,testAbsorption)
accuracy1;-data.frame(obs=testFat,pred=prediction)
defaultSummary(accuracy1)
plot(accuracy1)

```

```

#####
#Question 6.2
#####

```

```

library(AppliedPredictiveModeling)
library(MASS)
library(caret)
library(elasticnet)
library(lars)
library(pls)

```

```

data(permeability)

```

```

#####
# question 6.2(b)
#####
cat("Before Non-Zero Variance, number of predictors in fingerprints is 1107: "n")
print(str(fingerprints))
cat("n")

```

```

cat("After Non-Zero Variance, number of predictors in fingerprints is 388: "n")
NZVfingerprints ; nearZeroVar(fingerprints)
noNZVfingerprints ; fingerprints[, -NZVfingerprints]
print(str(noNZVfingerprints))
cat("n")

```

```

#####
#####

```

```

#####
# question 6.2(c)
#####

```

```

# stratified random sample splitting with 75% training and 25% testing

```

```

set.seed(12345)
trainingRows = createDataPartition(permeability, p = .75, list= FALSE)

```

```

trainFingerprints ; noNZVfingerprints[trainingRows,]

```

```

trainPermeability <- permeability[trainingRows,]

testFingerprints <- nonZVfingerprints[-trainingRows,]
testPermeability <- permeability[-trainingRows,]

set.seed(12345)

ctrl <- trainControl(method = "repeatedcv", repeats=5, number = 4)

# PLS Model
permeabilityPLS <- train(x = trainFingerprints , y = trainPermeability, method = "pls", tuneGrid = expand.grid(
print(permeabilityPLS)
plot(permeabilityPLS, metric ="Rsquared", main = "PLS Tuning Parameter for Permeability Data")

# PLS prediction
prediction<-predict(permeabilityPLS,testFingerprints)
accuracy <- data.frame(obs=testPermeability, pred=prediction)
defaultSummary(accuracy)
plot(accuracy)

cat("n")

# Ridge Regression Method
permeabilityRg <- train(x = trainFingerprints , y = trainPermeability, method = "ridge",
trControl = ctrl,
tuneGrid = expand.grid(lambda = seq(0.08,0.25,length=15)))

prediction<-predict(permeabilityRg,testFingerprints)
accuracy <- data.frame(obs=testPermeability, pred=prediction)
defaultSummary(accuracy)
plot(accuracy)

cat("n")

#
# Elastic Net Method
PermeabilityEls <- train(x = trainFingerprints , y = trainPermeability, method = "enet",
trControl = ctrl,
tuneGrid = expand.grid(lambda = c(.001,.01,.1),
fraction = seq(0.05,1,length=20)))

print(permeabilityEls)
plot(permeabilityEls)
cat("n")

# Elastic Net Method Prediction
prediction<-predict(PermeabilityEls,testFingerprints)
accuracy <- data.frame(obs=testPermeability, pred=prediction)
defaultSummary(accuracy)
plot(accuracy)

cat("n")

# Lasso Regression Method
permeabilityLasso <- train(x = trainFingerprints , y = trainPermeability, method = "lasso",
trControl = ctrl,
tuneGrid = expand.grid(fraction = seq(0.0005,0.005,length=20)))

print(permeabilityLasso)
plot(permeabilityLasso)
cat("n")
#

```

```
# # Lasso Model Prediction
prediction;-predict(PermeabilityLasso,testFingerprints)
accuracy ;- data.frame(obs=testPermeability, pred=prediction)
defaultSummary(accuracy)
plot(accuracy)

#
# #####
# #####
#
```


References:

1. Applied Predictive Modeling : @authors Max Kuhn. Kjell Johnson
2. <https://archive.ics.uci.edu/ml/index.php>